

Selection-track composition leaderboard — best: Raw + all subj. encodings (cosine)
aggregated across model sets; x = in-silico model-recovery AUC

